

The Network Contributor Rewards Model

Nihar Shah

DoubleZero Foundation

July 17, 2025

Outline

1. Running the Model in Practice
2. Motivating Shapley Values
3. Simple Example
4. Adjustments and Complications

Outline

1. Running the Model in Practice
2. Motivating Shapley Values
3. Simple Example
4. Adjustments and Complications

Running the Model in Practice

- ▶ Full model in `network_shapley.py` in repo
<https://github.com/doublezerofoundation/network-shapley>
- ▶ Readme, long PDF, and example inputs are in repo
- ▶ Production version of the code being developed in Rust in a linked repo
- ▶ We'll walk through a minimal working example here

Minimal Working Example

```
import pandas as pd
from network_shapley import network_shapley

private_links = pd.read_csv("private_links.csv")
devices       = pd.read_csv("devices.csv")
public_links  = pd.read_csv("public_links.csv")
demand1       = pd.read_csv("demand1.csv")

result = network_shapley(
    private_links      = private_links,
    devices            = devices,
    demand             = demand1,
    public_links       = public_links,
    operator_uptime    = 0.98, # optional
    contiguity_bonus   = 5.0,  # optional
    demand_multiplier  = 1.2   # optional
)

print(result)
```

Required Inputs

- ▶ **Private links table:** one row per private link (bidirectional)
(Device1, Device2, Latency, Bandwidth, Uptime, Shared)
- ▶ **Device table:** one row per device
(Device, Edge, Operator)
- ▶ **Public links table:** one row per public link (bidirectional)
(City1, City2, Latency)
- ▶ **Demand matrix:** one row per traffic flow
(Start, End, Receivers, Traffic, Priority, Type, Multicast)

Optional Inputs

- ▶ **Operator uptime:** probability that any given operator stays through epoch
- ▶ **Contiguity bonus:** latency penalty applied to hybrid private-public routing, avoided through using contiguous routing
- ▶ **Demand multiplier:** scalar to adjust demand up

Minimal Working Example

```
import pandas as pd
from network_shapley import network_shapley

private_links = pd.read_csv("private_links.csv")
devices       = pd.read_csv("devices.csv")
public_links  = pd.read_csv("public_links.csv")
demand1       = pd.read_csv("demand1.csv")

result = network_shapley(
    private_links      = private_links,
    devices            = devices,
    demand             = demand1,
    public_links       = public_links,
    operator_uptime    = 0.98, # optional
    contiguity_bonus   = 5.0,  # optional
    demand_multiplier  = 1.2   # optional
)

print(result)
```

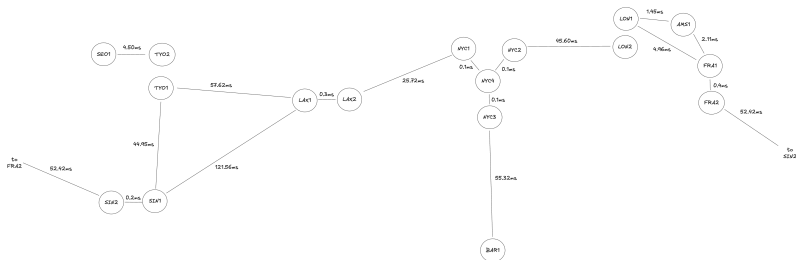

Output

- ▶ Returns a `pandas.DataFrame` with columns:

Operator	Name
Value	Raw Shapley value
Percent	$\frac{\text{Value}}{\sum \text{Value}}$

- ▶ Links earn rewards based on ability to cut latency and where traffic demand is
- ▶ So expect to see large changes depending on where leader is

Simulated Scenario



Simulated Scenario

Operator Name	Value #1	Percent #1	Value #2	Percent #2
Alpha	21.5370	2.80%	2.0154	0.16%
Beta	10.6595	1.03%	187.1199	15.01%
Delta	13.5257	1.31%	111.6727	8.96%
Epsilon	0.0407	0.00%	88.5022	7.10%
Gamma	487.1094	47.01%	23.0343	1.85%
Kappa	0.0603	0.01%	10.6422	0.85%
Theta	503.1153	48.55%	333.5523	26.76%
Zeta	0.1393	0.01%	490.0349	39.13%


Outline

1. Running the Model in Practice
2. Motivating Shapley Values
3. Simple Example
4. Adjustments and Complications


Alternate Idea: Carried Traffic Model

- ▶ Carried-traffic model pays only for traffic on each link... seems simple, right?
- ▶ But... it is insufficiently discriminating and cannot be kept simple for long
- ▶ Shapley values are much more discriminating and robust
- ▶ They pay out for the *marginal* contribution to a common **value function**

Illustrative Example



Illustrative Example




Shapley Values

- ▶ Carried-traffic model pays only for traffic on each link... seems simple, right?
- ▶ But... it is insufficiently discriminating and cannot be kept simple for long
- ▶ Shapley values are much more discriminating and robust
- ▶ They pay out for the *marginal* contribution to a common **value function**

Outline

1. Running the Model in Practice
2. Motivating Shapley Values
3. Simple Example
4. Adjustments and Complications

Scenario



Value Function

$$V = - \sum_i t_i l_i$$

- ▶ Manifestation of IBRL
- ▶ t_i : traffic flow i
- ▶ l_i : latency incurred by flow
- ▶ Negative sign means we want to maximize this term

Goal: compute this for every coalition of possible contributors

Linear Program Formulation

- ▶ To get value for a given set of link operators...

$$\begin{aligned} & \text{minimize} && c^\top x \\ & \text{subject to} && A_{eq} x = b_{eq}, \\ & && A_{ub} x \leq b_{ub}, \\ & && x \geq 0 \end{aligned}$$

- ▶ Decision variables x : traffic routed across each directed edge for each traffic type
- ▶ Solving yields coalition-specific value $V(C) = -c^\top x^*$

Private and Public Link Inputs

Start	End	Latency (ms)	Operator
<i>A</i>	<i>B</i>	10	Red
<i>A</i>	<i>B</i>	25	Green
<i>B</i>	<i>C</i>	20	Blue
<i>A</i>	<i>C</i>	25	Orange
<i>A</i>	<i>B</i>	30	Public
<i>B</i>	<i>C</i>	30	Public
<i>A</i>	<i>C</i>	30	Public

- Each bidirectional link will be split into two directed edges by the model

Flow-Conservation Matrix

$$\tilde{A}_{eq} = \begin{bmatrix} r & g & b & o & r' & g' & b' & o' & p_1 & p_2 & p_3 & p'_1 & p'_2 & p'_3 \\ 1 & 1 & 0 & -1 & -1 & -1 & 0 & 1 & 1 & 0 & -1 & -1 & 0 & 1 \\ -1 & -1 & 1 & 0 & 1 & 1 & -1 & 0 & -1 & 1 & 0 & 1 & -1 & 0 \\ 0 & 0 & -1 & 1 & 0 & 0 & 1 & -1 & 0 & -1 & 1 & 0 & 1 & -1 \end{bmatrix}$$

- ▶ Rows are nodes and columns are directed edges
- ▶ +1 if traffic leaves the node, -1 if it enters the node
- ▶ Need different flow matrices for different traffic types

$$A_{eq} = \text{diag}(\tilde{A}_{eq}, \tilde{A}_{eq}, \tilde{A}_{eq})$$

Demand Vectors

$$\tilde{b}_{eq}^{(1)} = \begin{bmatrix} 1 \\ -1 \\ 0 \end{bmatrix}, \quad \tilde{b}_{eq}^{(2)} = \begin{bmatrix} 0 \\ 1 \\ -1 \end{bmatrix}, \quad \tilde{b}_{eq}^{(3)} = \begin{bmatrix} -1 \\ 0 \\ 1 \end{bmatrix}$$

- ▶ Again, need different flow matrices for different traffic types

$$b_{eq} = \begin{bmatrix} \tilde{b}_{eq}^{(1)} \\ \tilde{b}_{eq}^{(2)} \\ \tilde{b}_{eq}^{(3)} \end{bmatrix}$$

Bandwidth Constraints

- ▶ Suppose all links but Green have 5 gigabits of capacity; and the Green has 2 gigabits
- ▶ Public edges have effectively infinite capacity so there is no constraint imposed on them

$$\tilde{A}_{ub} = [I_{8 \times 8} \ 0_{8 \times 6}]$$

$$\tilde{b}_{ub} = [5 \ 2 \ 5 \ 5 \ 5 \ 2 \ 5 \ 5]^T$$

- ▶ The matrices are stacked horizontally here, because all traffic types draw from the same bandwidth constraint

$$A_{ub} = [\tilde{A}_{ub} \ \tilde{A}_{ub} \ \tilde{A}_{ub}]$$

Cost Vector

- ▶ Objective minimizes latency, i.e. cost, across traffic types

$$\tilde{c} = [10 \quad 25 \quad 20 \quad 25 \quad 10 \quad 25 \quad 20 \quad 25 \quad 30 \quad 30 \quad 30 \quad 30 \quad 30 \quad 30]$$

- ▶ Latency is replicated across the three traffic types

$$c = [\tilde{c} \quad \tilde{c} \quad \tilde{c}]$$

- ▶ This defines all the components needed for the linear program

Value Function Across Coalitions

Coalition	I_{AB}	I_{BC}	I_{CA}	$V(C)$
No Operators	30	30	30	-90
Red	10	30	30	-70
Green	25	30	30	-85
Blue	30	20	30	-80
Orange	30	30	25	-85
Red, Blue	10	20	30	-60
Red, Orange	10	30	25	-65
Green, Blue	25	20	30	-75
Green, Orange	25	30	25	-80
Blue, Orange	30	20	25	-75
Red, Green	10	30	30	-70
Red, Blue, Orange	10	20	25	-55
Green, Blue, Orange	25	20	25	-70
Red, Green, Blue	10	20	30	-60
Red, Green, Orange	10	30	25	-65
All Operators	10	20	25	-55

Green Link's Marginal Contribution

With G	Without G	ΔV	Weight
$V(G, R, B, O)$	$V(R, B, O)$	$(-55) - (-55) = 0$	0.250
$V(G, R, B)$	$V(R, B)$	$(-60) - (-60) = 0$	0.083
$V(G, R, O)$	$V(R, O)$	$(-65) - (-65) = 0$	0.083
$V(G, B, O)$	$V(B, O)$	$(-70) - (-75) = 5$	0.083
$V(G, R)$	$V(R)$	$(-70) - (-70) = 0$	0.083
$V(G, O)$	$V(O)$	$(-80) - (-85) = 5$	0.083
$V(G, B)$	$V(B)$	$(-75) - (-80) = 5$	0.083
$V(G)$	$V(\text{none})$	$(-85) - (-90) = 5$	0.250
Sum			2.50

- ▶ Summing weighted differences yields Green's Shapley value = 2.5
- ▶ Red = 17.5 (50%), Blue = 10 (29%), Orange = 5 (14%), Green = 2.5 (7%).

Outline

1. Running the Model in Practice
2. Motivating Shapley Values
3. Simple Example
4. Adjustments and Complications

Operator Redundancy

- ▶ Operators may withdraw for commercial, operational, or regulatory reasons
- ▶ This methodology would over-reward fragile topologies
- ▶ Solution: pay contributors according to the *expected* value of a network and account for probability p of withdrawals
- ▶ This preemptively rewards links with insurance value

$$V = -\mathbb{E} \left(\sum_i t_i l_i \right)$$

Operator Redundancy

- Formally, turn coalition-specific values into expected values:

$$\mathbb{E}[V(C)] = \sum_{S \subseteq C} p^{|S|} (1-p)^{|C|-|S|} V(S)$$


- In practice, this is a computational bottleneck so we use a one-pass algorithm (described in the manual)

$$\mathbb{E}(V) = CB_V$$

DZXs, Edges, and Contiguity

- ▶ This methodology would ideally incentivize intra-city connections (“DZXs”) and network edges too
- ▶ But edges are outside the methodology; and the scope for DZXs to improve over the public internet is small given short distances
- ▶ Solution is to add a fixed latency penalty to non-contiguous (a.k.a. hybrid) routes, i.e. routes that need to use the public internet to transit through a city or off the network to users
- ▶ This encourages end-to-end deployments, but still admits non-contiguous networks as needed

DZXs, Edges, and Contiguity



- ▶ Orange public link is charged some penalty
- ▶ Contributors only unlock the full improvements for LA-London traffic by building the green private link

Measuring Latency

- ▶ Real-world latencies are hard to measure as a single number, because they fluctuate with congestion and size
- ▶ Solution: measure latency distributions over sampled intervals and non-trivial flow sizes
- ▶ Use an upper percentile (e.g. p95) as the effective latency measurement in the value function
- ▶ Brings notion of both speed and jitter into rewards calculus without changing core structure

Scaling Demand

- ▶ There may be different priorities of traffic (e.g. stake weight of different receivers); so the value function is extended:

$$V = -\mathbb{E} \left(\sum_i p_i t_i l_i \right)$$

- ▶ Also, potentially scale current demand by $\gamma > 1$ to anticipate future growth and expose chokepoints
- ▶ Gives incentives to invest in these before demand catches up
- ▶ Gradually adjust $\gamma \rightarrow 1$ as the network matures

Multicast

- ▶ Multicast is the ability for the network, rather than the sender, to duplicate packets
- ▶ This conserves bandwidth and allows for higher link utilization
- ▶ The methodology already extends to this case, for multicast-enabled demand flows
- ▶ All links, except links that move traffic from the network to the public internet, automatically support multicast in DoubleZero

Conclusion

- ▶ Check out the repo, readme, and manual
- ▶ There is a video of an *earlier* version of this deck and methodology at <https://youtu.be/K1Ni-k51sMw>
- ▶ Contact me at nihar@doublezero.us with questions
- ▶ On the horizon...
 1. The primary focus is turning this into production code for launch
 2. As part of that, there is a strong focus on measuring quantities and setting parameters accurately
 3. In the medium term, it would be nice to design a web interface